Zilong Zheng

Contact Information

zlzheng.cs@gmail.com
https://zilongzheng.github.io
https://scholar.google.com/citations?user=9sDx70IAAAAJ

Education

University of California, Los Angeles — **Los Angeles, CA**
Ph.D. in Computer Science — Jan. 2018 - June. 2021
M.S. in Computer Science — Sept. 2016 - Dec. 2017

University of Minnesota, Twin Cities — **Minneapolis, MN**
B.A. in Computer Science, **Distinction** — Jan. 2015 - June. 2016

University of Electronic Science and Technology of China — **Chengdu, China**
B.E. in Microelectronic Technology, Rank: **1/164** — Sept. 2012 - Jan. 2015

Research Interests

Faithful Reasoning: Developing reasoning language models that adhere to logic, symbolic rules, factual knowledge, and commonsense.
Human-AI Interactions & Agents: Building safe, self-explainable, and socially aware language agents aligned with human intent, ethics, social norms, and values.
High-dimensional Generative Modeling: Modeling high-dimensional data via energy-based models and diffusion processes for controllable generation.

Professional Experience

Beijing Institute for General Artificial Intelligence (BIGAI) — Jun. 2021 - Present
Principal Researcher of Agentic AI | Ex-Director of Research Department
Research Director of Natural Language and Conversational AI (NLCo) Lab
https://github.com/bigai-nlco

- Lead NLP/LLM and agentic-AI research; published 60+ papers in top-tier venues; developed faithful reasoning models, LLM acceleration methods, and interactive conversational/agentic systems; statistical generative modeling; mentoring 10+ PhD students from top institutions (PKU, THU, SJTU, ZJU, USTC, etc.).
- Lead **TongAgents**: Built large-model agentic training/inference systems and efficiency techniques (e.g., long-context decoding acceleration; native sparse/parallel inference). Achieving SOTA performance (surpassing GPT-o1 and DeepSeek-R1) and $\times$ 3 speedup across all domain-finetuned models on project TongLaw.
- Principal Investigators (PIs)/Co-PIs of over 5M funding on national projects (NSFC General Program + key R&D), leading multimodal, alignment, and agentic systems.
- Managed computing resources for the research department: procured and operated A100 clusters, overseeing deployment, hosting, and Slurm scheduling; responsible for hiring and mentoring researchers, engineers, and students.

Center for Vision, Cognition, Learning and Autonomy (VCLA) — Jan. 2017 - June. 2021
Graduate Student Researcher (Supervisor: Prof. Song-Chun Zhu)

- Thesis: Multimodal Conversation Modeling via Neural Perception, Structure Learning, and Communication

Baidu Research, Cognitive Computing Laboratory — June 2020 - Sept. 2020
Research Scientist Intern (Supervisor: Prof. Ping Li)

- Multimodal Representation Learning (CVPR'21 Oral).

Google, Mountain View, CA — June. 2017 - Sept. 2017
Software Engineer Intern

- Google Assistant Backend: Implementing acoustic controlling feature using C++

Selected Projects & Publications

Native Parallel Reasoner (NPR) July. 2025 - Nov. 2025
No. 1 Paper of the Day on Huggingface

- Teacher-free framework for **native parallel reasoning** via progressive self-distillation (ZERO → BETA → RL) and parallel-aware policy optimization.
- Built an SGLang-based rollout engine with KV reuse and adaptive branching; **up to +24.5%** and **4.6**× faster inference across 8 benchmarks.

Absolute Zero Reasoner (AZR) NeurIPS'25 Spotlight Feb. 2025 - May. 2025
Github Stars: 1.8K+, X Views: 495K+, No. 1 Paper of the Day on Huggingface, AlphaXiv Trending No. 1. Featured news on WIRED.

- Zero-data self-play: the model **generates tasks and solves them** in a verifiable Python environment (deduction/abduction/induction).
- Demonstrated strong scaling and transfer (e.g., **+15.2** math from code-only self-play); open-sourced training/eval pipeline.

TokenSwift: Ultra-Long Decoding Speedup ICML'25 Sept. 2024 - Jan. 2025

- Lossless generation acceleration up to **100K-token** with speculative decoding, featuring multi-token generation, token reuse, and dynamic KV-cache management.
- Achieved ∼**3**× speedup on any Language Models (e.g., Qwen2.5-32B), reducing generation time from 5 hours to 90 minutes; introduced contextual penalties to reduce long-horizon repetition.

OmniMMI: Streaming OmniLLM CVPR'25 | ICCV'25 Sept. 2024 - Dec. 2024

- Proposed first-ever multimodal **multiplexing** architecture with parallel decoding to **see/listen while generating**, enabling low-latency streaming interaction.
- Developed an efficient framework for long-form video understanding, achieving SOTA performance across four video understanding and egocentric planning tasks, with strong scalability (up to 8× video length) and excelling in fine-grained frame retrieval.

TongVerse: Embodied Testbed IROS'25 Oral Apr. 2023 - Jan. 2024
Foundation embodied platform for CRAIC 24,25.

- Open humanoid learning/evaluation suite based on Nvidia Isaac Sim with reproducible setups for perception–planning–control pipelines.
- Designed task suites + infrastructure to study policy learning under complex dynamics and diverse embodied objectives.

Bidirectional Value Alignment Science Robotics'22 Dec. 2020 - July. 2022
Featured as headline news on Science and Science Robotics, July 2022. Ref. XinhuaNet, ScienceNet, TechXplore Report.

- Proposed bidirectional alignment where robots infer human values from in-situ feedback and communicate decisions via explanations.
- Validated interactive alignment with user studies; emphasizes robustness in human–AI co-adaptive systems.

Energy-Based Modeling (EBM) + MCMC Teaching Dec. 2017 - May. 2021
CVPR'21 Oral | AAAI'21 | TPAMI'21 | TPAMI'20 | CVPR'19 Oral | CVPR'18 Oral

- Developed energy-based generative modeling methods using **Langevin dynamics / MCMC** and cooperative training (fast initializer + slow solver) for stable high-dimensional sampling.
- Contributions span patchwise EBMs, 3D descriptor networks, and alternating MCMC teaching for unsupervised translation and conditional generation.

Professional Activities

Conference/Journal Area Chairs/Reviewers

- ❑ Senior Area Chair, Empirical Methods in Natural Language Processing (EMNLP), 2025
- ❑ Senior Area Chair, Association for Computational Linguistics (ACL), 2026
- ❑ Area Chair, International Conference on Learning Representations (ICLR), 2025.
- ❑ Area Chair, International Conference on Machine Learning (ICML), 2026.
- ❑ Journal Reviewer, International Journal on Computer Vision (IJCV), NeuroComputing, Pattern Recognition, Engineering

Workshop Organizers

- ❑ Primary Host, The 1st Joint Workshop on Large Language Models and Structure Modeling (XLLM), ACL 2025
- ❑ Host, General Embodied Intelligence in Real and Virtual Open Worlds, China Embodied AI Conference (CEAI), 2024, 2025
- ❑ Host, Open-World Agents: Synergizing Reasoning and Decision-Making in Open-World Environments (OWA-2024), NeurIPS 2024
- ❑ Primary Host, Unimodal and Multimodal Induction of Linguistic Structures (UM-IoS), EMNLP 2021

Publications

(*: equal contributions, †: corresponding author)

Faithful Reasoning and Language Modeling

[60] T. Wu, Y. Liu, J. Bai, Z. Jia, S. Zhang, Z. Lin, Y. Wang, S.-C. Zhu and **Z. Zheng**†. Native Parallel Reasoner: Reasoning in Parallelism via Self-Distilled Reinforcement Learning. 2025.

[59] Y. Liu, J. Li and **Z. Zheng**†. RuleReasoner: Reinforced Rule-based Reasoning via Domain-aware Dynamic Sampling. 2025.

[58] H. Li, C. Li, T. Wu, X. Zhu, Y. Wang, Z. Yu, E. H. Jiang, S.-C. Zhu, Z. Jia, Y. N. Wu† and **Z. Zheng**†. Seek in the Dark: Reasoning via Test-Time Instance-Level Policy Gradient in Latent Space. 2025.

[57] A. Zhao, Y. Wu, Y. Yue, T. Wu, Q. Xu, Y. Yue, M. Lin, S. Wang, Q. Wu, **Z. Zheng**† and G. Huang†. Absolute Zero: Reinforced Self-play Reasoning with Zero Data. In *the 39th Conference on Neural Information Processing Systems (NeurIPS)*, 2025. **[Spotlight]**

[56] X. Qin, J.Bai, J. Li, Z. Jia†, and **Z. Zheng**†. Reinforced Query Reasoners for Reasoning-intensive Retrieval Tasks. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.

[55] J. Bai, M. Tong, Y. Liu, Z. Jia†, and **Z. Zheng**†. Understanding and Leveraging the Expert Specialization of Context Faithfulness in Mixture-of-Experts LLMs. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2025.

[54] Z. Lin*, H. Wu*, S. Wang, K. Tu†, **Z. Zheng**†, and Z. Jia†. Look Both Ways and No Sink: Converting LLMs into Text Encoders without Training. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.

[53] J. Li, X. Dong, Y. Liu, Z. Yang, Q. Wang, X. Wang, S.-C. Zhu, Z. Jia†, and **Z. Zheng**†. ReflectEvo: Improving Meta Introspection of Small LLMs by Learning Self-Reflection. In *Findings of the Association for Computational Linguistics: ACL*, 2025.

[52] T. Wu*, J. Shen*, Z. Jia, Y. Wang and **Z. Zheng**†. TokenSwift: Lossless Acceleration of Ultra Long Sequence Generation. In *Proceedings of the Forty-second International Conference on Machine Learning (ICML)*, 2025.

[51] X. Zhu, D. Cheng, H. Li, K. Zhang, E. Hua, X. Lv, N. Ding, Z. Lin†, **Z. Zheng**† and B. Zhou†. How to Synthesize Text Data without Model Collapse? In *Proceedings of the Forty-second International Conference on Machine Learning (ICML)*, 2025.

[50] S. Qi†, B. Yang, K. Jiang, X. Wang, J. Li, Y. Zhong, Y. Yang and **Z. Zheng**†. In-Context Editing: Learning Knowledge from Self-Induced Distributions. In *Proceedings of the Thirteenth International Conference on Learning Representations (ICLR)*, 2025.

[49] A. Zhao, Q. Xu, M. Lin, S. Wang, Y.-J. Liu, **Z. Zheng**† and G. Huang†. DiveR-CT: Diversity-enhanced Red Teaming Large Language Model Assistants with Relaxing Constraints. In *Proceedings of the AAAI Conference on Artificial Intelligence (AAAI)*, 2025. **[Oral]**

[48] T. Wu, Y. Zhao and **Z. Zheng**†. An Efficient Recipe for Long Context Extension via Middle-Focused Positional Encoding. In *the 38th Conference on Neural Information Processing Systems (NeurIPS)*, 2024.

[47] X. Tang, J. Li, Y. Liang, M. Zhang and **Z. Zheng**†. Mars: Situated Inductive Reasoning in an Open-World Environment. In *the 38th Conference on Neural Information Processing Systems (NeurIPS) Track on Datasets and Benchmarks*, 2024.

[46] Z. Lin, Q. Wang, Z. Jia† and **Z. Zheng**†. Varying Sentence Representations via Condition-Specified Routers. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.

[45] Y. Wang, A. Yuille, Z. Li† and **Z. Zheng**†. ExoViP: Step-by-step Verification and Exploration with Exoskeleton Modules for Compositional Visual Reasoning. In *The First Conference on Language Modeling (CoLM)*, 2024.

[44] J. Li, M. Wang, **Z. Zheng**† and M. Zhang†. LooGLE: Can Long-Context Language Models Understand Long Contexts? In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.

[43] Z. Jia, J. Li, S. Zhang and **Z. Zheng**†. Combining Supervised Learning and Reinforcement Learning for Multi-Label Classification Tasks with Partial Labels. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (ACL)*, 2024.

[42] J. Li*, Z. Jia* and **Z. Zheng**†. Semi-automatic Data Enhancement for Document-Level Relation Extraction with Distant Supervision from Large Language Models. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.

[41] Z. Jia, Z. Yan, W. Han, **Z. Zheng**† and K. Tu †. Modeling Instance Interactions for Joint Information Extraction with Neural High-Order Conditional Random Field. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.

[40] Y. Wang, J. Wang, D. Zhao† and **Z. Zheng**†. Shuō Wén Jiě Zì: Rethinking Dictionaries and Glyphs for Chinese Language Pre-training. *Findings of the association for computational linguistics: ACL*, 2023.

[39] L. Zhang, Z. Jia, W. Han, **Z. Zheng** and K. Tu. SHARP: Search-Based Adversarial Attack for Structured Prediction. In *Findings of the Association for Computational Linguistics: NAACL*, 2022.

[38] Y. Li, **Z. Zheng**, W. Han and L. Zou. VGStore: A Multimodal Extension to SPARQL for Querying RDF Scene Graph. In *ISWC Poster & Demo Track*, 2022.

Human-AI Interactions and Alignments

[37] L. Yuan*, X. Gao*, **Z. Zheng***, M. Edmonds†, Y. N. Wu, F. Rossano, H. Lu†, Y. Zhu† and S.-C. Zhu†. In situ bidirectional human-robot value alignment, *Science Robotics*, July 2022.

[36] X. Wang, Z. Jia, J. Li, Q. Liu and **Z. Zheng**†. Adaptive Preference Optimization with Uncertainty-aware Utility Anchor. In *Findings of the Association for Computational Linguistics: EMNLP*, 2025.

[35] H. Li, Z. Jiao, X. Liu†, H. Liu†, and **Z. Zheng**†. In-situ Value-aligned Human-Robot Interactions with Physical Constraints. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2025. **[Oral]**

[34] Y. Kang, J. Wang, Y. Li, M. Wang, W. Tu, Q. Wang, H. Li, T. Wu, X. Feng, F. Zhong, and **Z. Zheng**†. Are the Values of LLMs Structurally Aligned with Humans? A Causal Perspective. In *Findings of the Association for Computational Linguistics: ACL*, 2025.

[33] X. Zheng*, H. Lin*, K. He, Z. Wang, **Z. Zheng**† and Y. Liang†. MCU: An Evaluation Framework for Open-Ended Game Agents. In *Proceedings of the Forty-second International Conference on Machine Learning (ICML)*, 2025. **[Spotlight]**

[32] Y. Wang, Y. Wang, B. Chen, T. Wu, D. Zhao and **Z. Zheng**†. OmniMMI: A Comprehensive Multi-modal Interaction Benchmark in Streaming Video Contexts. In *Proceedings of the Computer Vision and Pattern Recognition Conference (CVPR)*, 2025.

[31] Z. Zhang, F. Bai, Q. Chen, C. Ma, M. Wang, H. Sun, **Z. Zheng**† and Y. Yang†. Amulet: ReAlignment During Test Time for Personalized Preference Adaptation of LLMs. In *Proceedings of the Thirteenth International Conference on Learning Representations (ICLR)*, 2025.

[30] S. Qiu, M. Liu, H. Li, S.-C. Zhu and **Z. Zheng**†. MindDial: Enhancing Conversational Agents with Theory-of-Mind for Common Ground Alignment and Negotiation. In *Proceedings of the 25th Annual Meeting of the Special Interest Group on Discourse and Dialogue (SIGDIAL)*, 2024. **[Oral]**

[29] R. Gong, Q. Huang, X. Ma, H. Vo, Z. Durante, Y. Noda, **Z. Zheng**, D. Terzopoulos, F.-F. Li and J. Gao. MindAgent: Emergent Gaming Interaction. In *Findings of the Association for Computational Linguistics: NAACL*, 2024.

[28] S. Wang, C. Liu, **Z. Zheng**†, S. Qi, S. Chen, Q. Yang, A. Zhao, S. Wang, S. Song and G. Huang†. Boosting LLM Agents with Recursive Contemplation for Effective Deception Handling. In *Findings of the Association for Computational Linguistics: ACL*, 2024.

[27] Z. Jia*, M. Wang*, B. Tong, S.-C. Zhu and **Z. Zheng**†. LangSuit·E: Controlling, Planning, and Interacting with Large Language Models in Embodied Text Environments. In *Findings of the Association for Computational Linguistics: ACL*, 2024.

[26] H. Li, S.-C. Zhu and **Z. Zheng**†. DiPlomat: A Dialogue Dataset for Situated Pragmatic Reasoning. In *the 37th Conference on Neural Information Processing Systems (NeurIPS) Track on Datasets and Benchmarks*, 2023.

[25] Y. Wang, **Z. Zheng**†, X. Zhao, J. Li, Y. Wang and D. Zhao†. VSTAR: A Video-grounded Dialogue Dataset for Situated Semantic Understanding with Scene and Topic Transitions. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.

[24] L. Fan*, S. Qiu*, **Z. Zheng**, T. Gao, S.-C. Zhu and Y. Zhu. Learning Triadic Belief Dynamics in Nonverbal Communication from Videos. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021. **[Oral]**

[23] **Z. Zheng**, S. Qiu, L. Fan, Y. Zhu and S.-C. Zhu. GRICE: A Grammar-based Dataset for Recovering Implicature and Conversational rEasoning. In *Findings of the Association for Computational Linguistics: ACL-IJCNLP*, 2021.

[22] T. Yuan, H. Liu, L. Fan, **Z. Zheng**, T. Gao, Y. Zhu and S.-C. Zhu. Joint Inference of States, Robot Knowledge, and Human (False-)Beliefs. In *IEEE International Conference on Robotics and Automation (ICRA)*, 2020.

[21] **Z. Zheng***, W. Wang*, S. Qi* and S.-C. Zhu. Reasoning Visual Dialogs with Structural and Partial Observations. In *Proceedings of IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2019. **[Oral]**

Multi-Modal Understanding and Generative Modeling

[20] Y. Wang*, Y. Song*, C. Xie, Y. Liu and **Z. Zheng**†. VideoLLaMB: Long-context Video Understanding with Recurrent Memory Bridges. In *Proceedings of International Conference on Computer Vision (ICCV)*, 2025.

[19] J. Zhu, Z. Du, H. Xu, F. Lan, **Z. Zheng**, B. Ma, S. Wang, and T. Zhang. Navi2Gaze: Leveraging Foundation Models for Navigation and Target Gazing. In *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2025. **[Oral]**

[18] Y. Peng*, Y. Ma*, M. Wang, Y. Wang, Y. Wang, C. Zhang, Y. Zhu†, and **Z. Zheng**†. Probing and Inducing Combinational Creativity in Vision-Language Models. In *Proceedings of the 47th Annual Meeting of the Cognitive Science Society (CogSci)*, 2025. **[Oral]**

[17] Y. Du*, K. Jiang*, Z. Gao, C. Shi, **Z. Zheng**†, S. Qi and Q. Li†. MMKE-Bench: A Multimodal Editing Benchmark for Diverse Visual Knowledge. In *Proceedings of the Thirteenth International Conference on Learning Representations (ICLR)*, 2025.

[16] Z. Wang, S. Cai, A. Liu, Y. Jin, J. Hou, B. Zhang. H. Lin, Z. He, **Z. Zheng**, Y. Yang, X. Ma† and Y. Liang†. JARVIS-1: Open-world Multi-task Agents with Memory-Augmented Multimodal Language Models. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 2024.

[15] Y. Wang, Y. Wang, P. Wu, J. Liang, D. Zhao, Y. Liu and **Z. Zheng**†. Efficient Temporal Extrapolation of Multimodal Large Language Models with Temporal Grounding Bridge for Long Video Understanding. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2024.

[14] J. Cui*, Z. Gong*, B. Jia*, S. Huang, **Z. Zheng**†, J. Ma† and Y. Zhu†. ProBio: A Protocol-guided Multimodal Dataset for Molecular Biology Lab. In *the 37th Conference on Neural Information Processing Systems (NeurIPS) Track on Datasets and Benchmarks*, 2023.

[13] X. Ma*, S. Yong*, **Z. Zheng**†, Q. Li, Y. Liang, S.-C. Zhu and S.Huang†. SQA3D: Situated Question Answering in 3D Scenes. In *the Eleventh International Conference on Learning Representations (ICLR)*, 2023. **Official challenge benchmark for workshop on 3D Scene Understanding for Vision, Graphics, and Robotics, CVPR 2025.**

[12] C. Lou*, W. Han, Y. Lin and **Z. Zheng***. Unsupervised Vision-Language Parsing: Seamlessly Bridging Visual Scene Graphs with Language Structures via Dependency Relationships. *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2022.

[11] B. Wan, W. Han, **Z. Zheng** and T. Tuytelaars. Unsupervised Vision-Language Grammar Induction with Shared Structure Modeling. In *The Tenth International Conference on Learning Representations (ICLR)*, 2022. **[Oral]**

[10] J. Zhang, J. Xie, **Z. Zheng** and N. Barnes. Energy-Based Generative Cooperative Saliency Prediction. In *Proceedings of the 36th AAAI Conference on Artificial Intelligence (AAAI)*, 2022. **[Oral]**

[9] J. Xie*, **Z. Zheng***, X. Fang, S.-C. Zhu and Y. N. Wu. Cooperative Training of Fast Thinking Initializer and Slow Thinking Solver for Multi-modal Conditional Learning. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 2021.

[8] **Z. Zheng**, J. Xie and P. Li. Patchwise Generative ConvNet: Training Energy-Based Models from a Single Natural Image for Internal Learning. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021. **[Oral]**

[7] J. Xie, Y. Xu, **Z. Zheng**, S.-C. Zhu and Y.N. Wu. Generative PointNet: Deep Energy-Based Learning on Unordered Point Sets for 3D Generation, Reconstruction and Classification. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, 2021.

[6] J. Xie*, **Z. Zheng***, X. Fang, S.-C. Zhu and Y.N. Wu. Learning Cycle-Consistent Cooperative Networks via Alternating MCMC Teaching for Unsupervised Cross-Domain Translation. In *Proceedings of the 35th AAAI Conference on Artificial Intelligence (AAAI)*, 2021.

[5] J. Xie, **Z. Zheng** and P. Li. Learning Energy-Based Model with Variational Auto-Encoder as Amortized Sampler. In *Proceedings of the 35th AAAI Conference on Artificial Intelligence (AAAI)*, 2021.

[4] J. Xie*, **Z. Zheng***, R. Gao, W. Wang, S.-C. Zhu and Y.N. Wu. Generative VoxelNet: Learning Energy-Based Models for 3D Shape Synthesis and Analysis. *IEEE Transactions on Pattern Analysis and Machine Intelligence (TPAMI)*, 2020.

[3] J. Xie*, R. Gao*, **Z. Zheng** and S.-C. Zhu. Motion-Based Generator Model: Unsupervised Disentanglement of Appearance, Trackable and Intrackable Motions in Dynamic Patterns. In *Proceedings of the 34th AAAI Conference on Artificial Intelligence (AAAI)*, 2020. **[Oral]**

[2] J. Xie*, R. Gao*, **Z. Zheng**, S.-C. Zhu and Y.N. Wu. Learning Dynamic Generator Model by Alternating Back-Propagation Through Time. In *Proceedings of the 33rd AAAI Conference on Artificial Intelligence (AAAI)*, 2019. **[Spotlight]**

[1] J. Xie*, **Z. Zheng***, R. Gao, W. Wang, S.-C. Zhu and Y.N. Wu. Learning Descriptor Networks for 3D Shape Synthesis and Analysis. *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018. **[Oral]**

Technical Skills

Languages: Python, JavaScript (ReactJS, TypeScript), C++, OCaml (teaching/TA experience), CUDA Kernel (2018-2019).
ML: PyTorch; Tensorflow; verl; SGLang; vllm.
Infra: Slurm; Kubernetes; Nvidia BCM.